

MultiValue Bookstore

Python Samples volume 1.

Python Sample 1

Sales Charting

Introduction

The Bookstore development team received the following user story into their backlog:

As a Sales Director
I want to visualize our sales performance across genres
So that we can focus on new promotions.

'Hey, I can do that', said Matthew, 'That is easy under Python. There are loads of charting libraries and most of them work with pandas'.

The team looked at him strangely. 'What's this got to do with pandas?' asked Sarah whose desk had pictures of black and white mammals.

'Not those pandas', Matthew explained. 'It's a standard Python library for manipulating data using in data science. You load it with data and then lots of other packages use it. That's the way the Python community works, we all help each other.'

'Wow', said James, 'we could do with some of that here'.

The others ignored him

'So what would you need?' asked Sarah more helpfully.

'I just need a way to grab the sales figures out of UniVerse', Matthew answered.

'Well that's easy on our side', said James. 'UniVerse has two enquiry languages: a native English- like one that is really friendly, and a version of SQL that is more powerful. You can use either of those to query the data. I can show you how to write a statement to retrieve that data.'

'How do I get the results? Pandas can run queries against SQL databases but it's kind of slow and needs setting up.'

James thought for a moment. 'Universe can produce XML from any query - Is that any good to you?'

Matthew nodded, happily. 'Sure, I can create an element tree and just iterate through that. Sounds like I can bring in UniVerse data really easily. Then I can give you the chart as an image.'

'Great', said James. 'Can we pair on this? I want to see how you do it!'

History

1. Proof of Concept

Matthew and James decided to pair on a simple stand-alone example first. James told Matthew how to write a simple UniVerse SQL query to produce XML.

```
SELECT GENRE,SUM(STOCK_LEVEL) FROM U2_BOOKS GROUP BY GENRE TOXML;
```

Matthew then used that to populate a pandas DataFrame and passed that into matplotlib to create a simple line chart. This was stored in a directory as a png.

See: **RUNPY books.pysrc plotstockgenre.py**

2. Final Version

Now James could see how to produce a chart in Python using UniVerse data, but they wanted to bind this into a stored procedure that could request a date range.

This needed a more complex UniVerse enquiry as it must normalize the multivalued lines of the sales orders selected to sum them.

UniVerse can run stand-alone Python routines and can call Python from within its business language, either as functions or by instantiating Python objects.

See: **U2.PLOT.GENRE.SALES, u2_plotGenreSales**

Python Sample 2

Recommendation Engine

Introduction

The Bookstore development team received the following user story into their backlog:

As a Sales Director
I want to improve our recommendations
So that we can realize more cross-selling opportunities.

After speaking with the director, the team decided to adopt a collaborative filtering approach similar to that used by other leading resellers (people who bought 'x' also bought 'y').

This would need to have the following features:

- A recommendation would be based on a publication and optionally a client.
- The recommendations would look at who had bought that publication, and what other publications those same customers had purchased.
- Repeat orders would not be counted.
- The recommendations would initially be ranked on popularity based on numbers. It is assumed (naively) that all publications have been available for the same length of time.
- Weightings would then be applied as follows:
 - For the same author a 3x weighting would be applied.
 - For the same genre a 2x weighting would be applied.
- Finally, if a client was specified for the recommendation, any books they had already purchased would be filtered from the recommendation.
- Only the top 10 recommendations would be returned along with their rankings, in the order of the most popular to the least.

History

1. Proof of Concept

books.bp u2_recommender_v1

In order to understand the requirements, a first attempt was made to build this using regular MultiValue techniques and the business language, by means of a Proof of Concept.

This created pairing records by examining the orders and writing combinations of titles purchased together to a U2_PAIRS file. A new U2_CLIENT_BOOKS file also kept track of which clients had purchased which titles, to speed up the lookups and to filter out purchases by the client to whom the recommendations were being shown. This could have been fulfilled using the existing U2_CLIENT_ORDERS but that would have slowed the retrieval of recommendations.

This proved workable in providing meaningful recommendations, but was slow in building and the team decided it could be improved upon.

See the test script: RECOMMENDER_1

2. First iteration

books.bp u2_recommender_v2

After meetings with the team to review the Proof of Concept, it was decided that the design of the U2_PAIRS and U2_CLIENT_BOOKS were potential bottlenecks to the process. Whilst the recommendations were correct, it would be better to cache these details in memory following the initial build to improve performance.

One team member mentioned that the U2 Dynamic Objects (UDOs) supported by the business language provided a good model for this, as these flexible JSON-style structures are ably suited to holding the key/value pairs that can match books with clients and vice versa. The proof of concept was rewritten to store details using UDOs in place of new data files.

This proved to be much more efficient, but there were concerns about the maintainability and extensibility of this approach.

See the test script: RECOMMENDER_2

3. Second iteration

```
books.pysrc u2_recommender_v3  
books.bp u2_recommender_v3
```

At this point, the team felt that they understood the problem domain and that the approach taken by the recommendation engine was correct. However, it would be better to develop this in a language that made the structure more apparent and that would be easier to extend in future.

Matthew, the new Python intern, suggested rewriting this as a Python routine. Python has the same data structures as UDOs but would be cleaner and simpler to maintain. It also had access to the same orders, books and client files as the business language.

The python routine would cache the recommendations along the same lines as the previous version. For this it would need to persist, and so Matthew suggested writing it as a service that could run in the background and that would be fed new sales orders as they were placed.

The team liked this approach, and another team member pointed out that the routine could run using the RUNPY command, as a UniVerse phantom process. All that was needed was a means for the two to communicate.

Matthew pointed out that Python has a fully serviceable simple http server, and recoded the recommendation engine to use this.

UniVerse and UniData also have built-in HTTP support, so the team were happy as they could switch from using low level socket calls in the Business Language to the simpler callHTTP requests.

Python Sample 3

Web

Introduction

Following the success of their first Python examples, the team were excited to learn more and to demonstrate their new Python skills. They realized that their existing Windows front end was not really showing this in its best light.

So the team asked the head of development to raise a user story on their behalf:

As a Development Manager
I want to create some web pages
So that we can demonstrate our application.

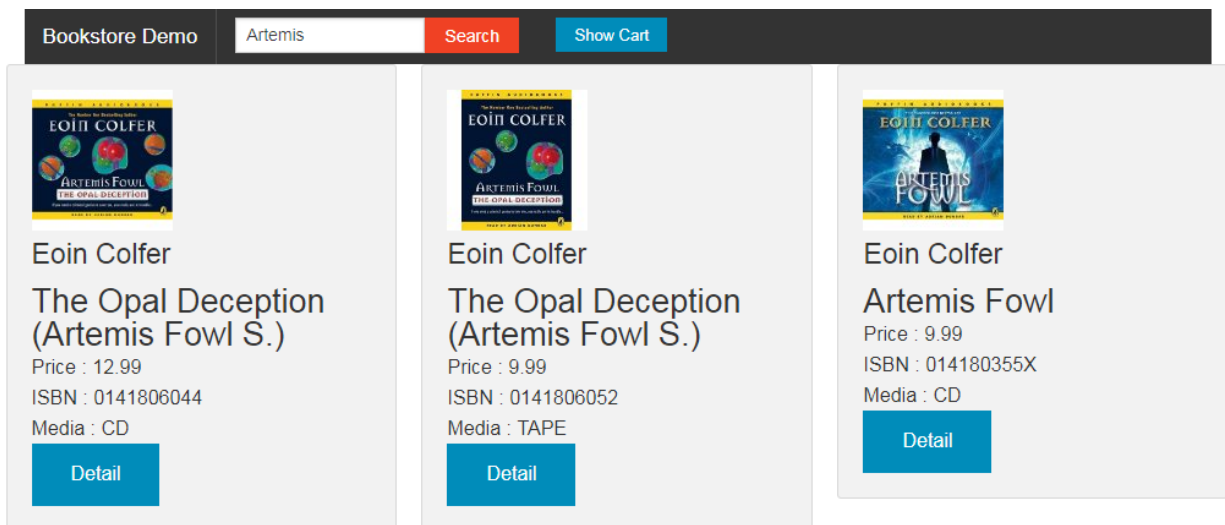
First Iteration

The new recommender has shown the team how easy it was to create an http listener using Python, and how well that could talk to UniVerse. Surely, they reasoned, this could be used to service more than the recommender? Could it be used to serve web pages and other web API requests?

This time, Matthew paired with Sarah to build a framework web server that could return pages using code found on the internet. They discovered the code was for Python 2, but in converting it to Python 3 Sarah learned a lot about how to work in a Pythonic manner.

Meanwhile, James worked with rest of the team to add in a set of methods that would interact with their UniVerse application. They decided on a convention of `/api/methodname` to call the method, and Matthew suggested that each method should be a separate class for ease of maintenance.

Before long they had a working web server that could model their bookstore as a Single Page App:

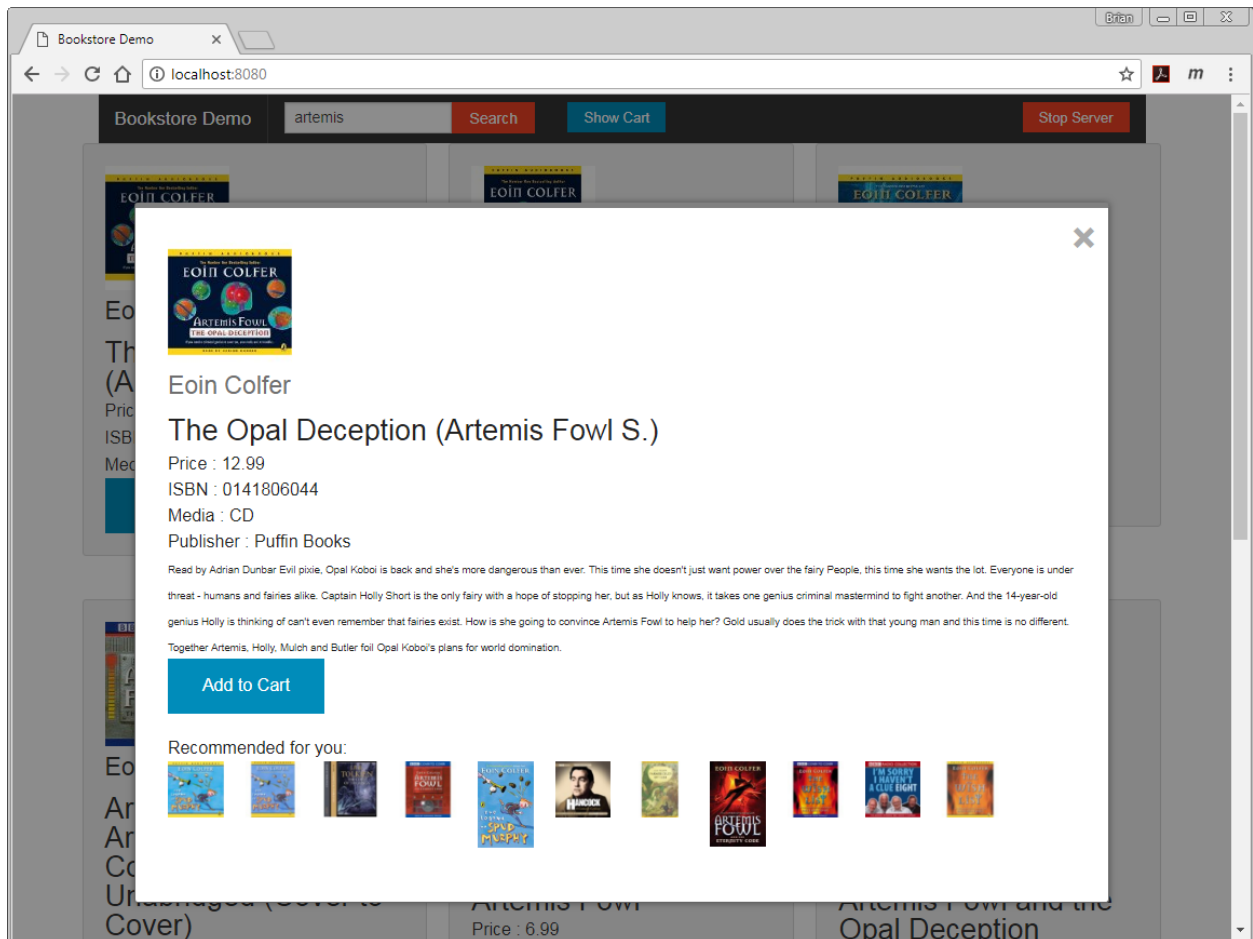


Second Iteration

books.pysrc webserver.py
books.pysrc u2_recommender_v3.py
web section of the repository

The team liked the new web server (though they argued long about the look and feel of the pages!).

But they did not like the fact that they now had two HTTP listeners to maintain. Could they be combined into one?



See:

UniVerse: start_webserver command